

Multicolumn Cache: aiming for both
high hit rate and low latency

Xiaodong Zhang

CSE 3421 additional notes for Memory Systems


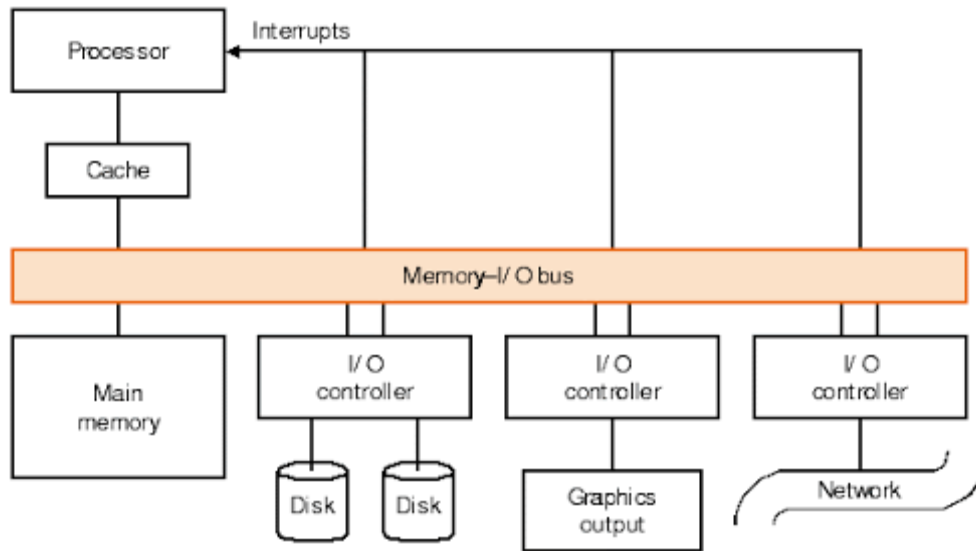
Basic concepts of hardware cache

- ❑ Cache design is **memory-address centric**
 - ❑ The default length is 32 bits, each memory address is unique and points to a byte (byte addressable)
 - ❑ Each memory address connects to the first byte of the content, and its default length is a word (4 Bytes or 32 bits)

- ❑ Cache is a small storage **inclusively** built in the CPU chip
 - ❑ It contains copies of the data from memory (and disks)
 - ❑ A cache-block storage can be multiple-words long
 - ❑ Besides data storage, each block needs tags V bit, and others

- ❑ For a given memory address, CPU knows where it is
 - ❑ By direct, set-associative or fully associative mapping
 - ❑ How do handle write hit, write miss, read hit, and read miss?
 - ❑ How to address the conflicts between **high hit rates** and **low latency**?

Connections between CPU-Caches and DRAM memory



Three steps for a cache miss:

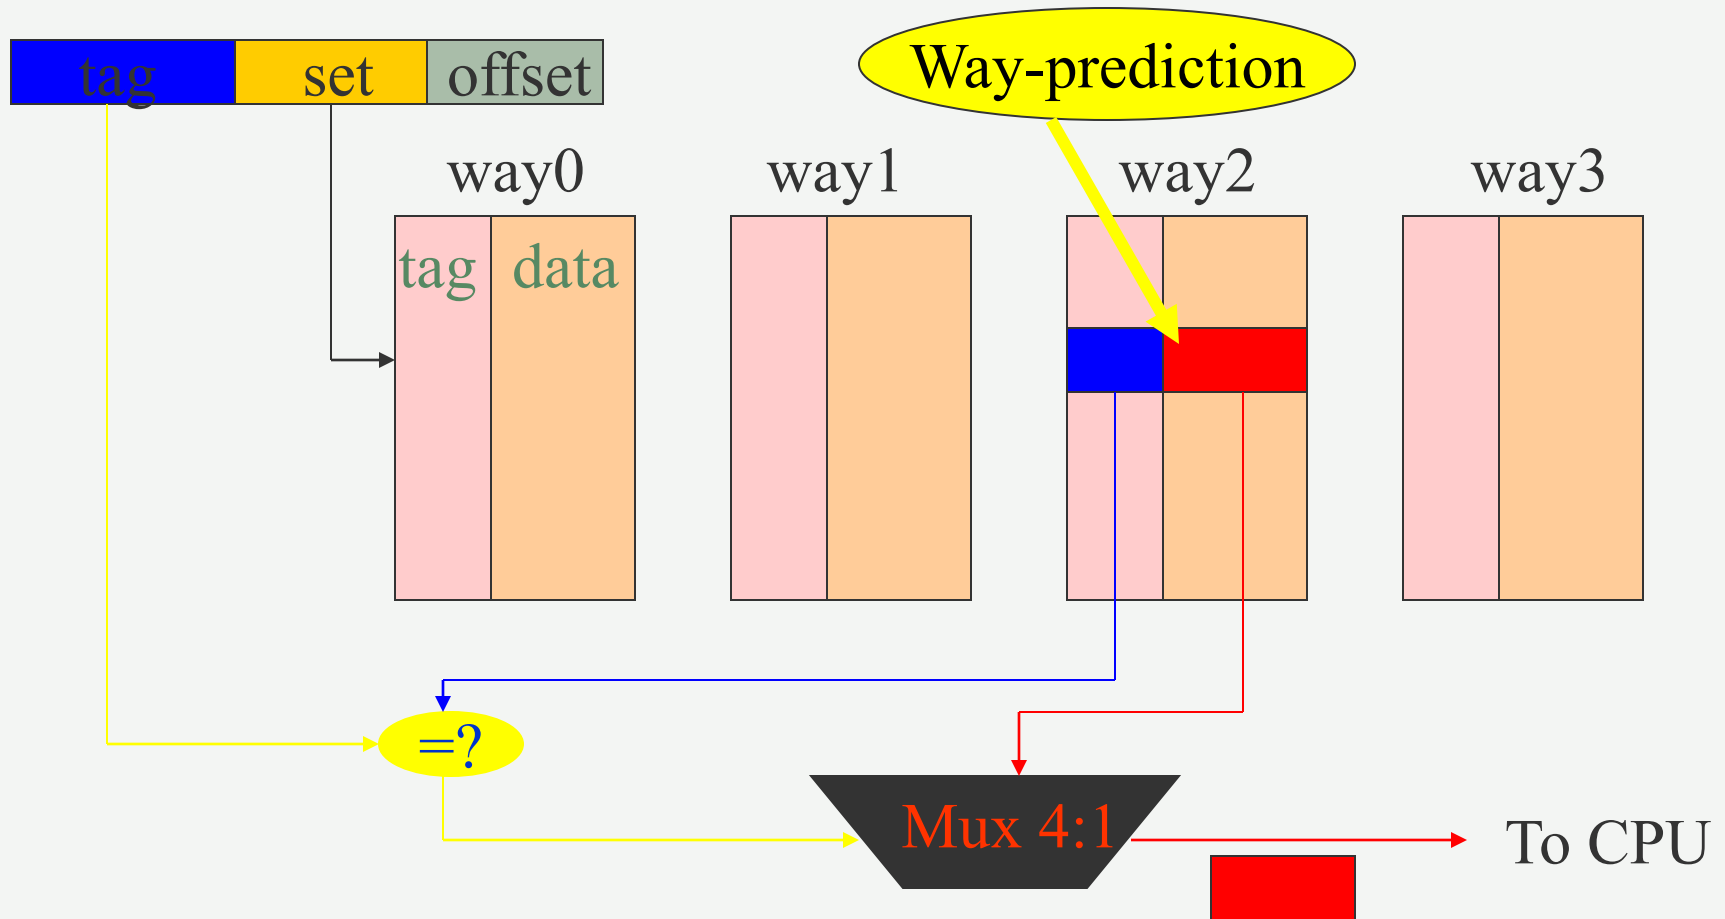
- ❑ Sending a memory address (bus)
- ❑ Loading data from memory (no bus)
- ❑ Transferring data to CPU (bus)

Trade-offs between High Hit Ratios and Low Latency

- Set-associative cache achieves high hit-ratios: 30% higher than that of direct-mapped cache.
- But it suffers high access times due to
 - Multiplexing logic delay during the selection.
 - Tag checking, selection, and data dispatching are sequential.
- Direct-mapped cache loads data and checks tag in parallel: minimizing the access time.
- High power consumption is another concern for set associative cache
- Can we get both high hit ratios and low latency in low power?

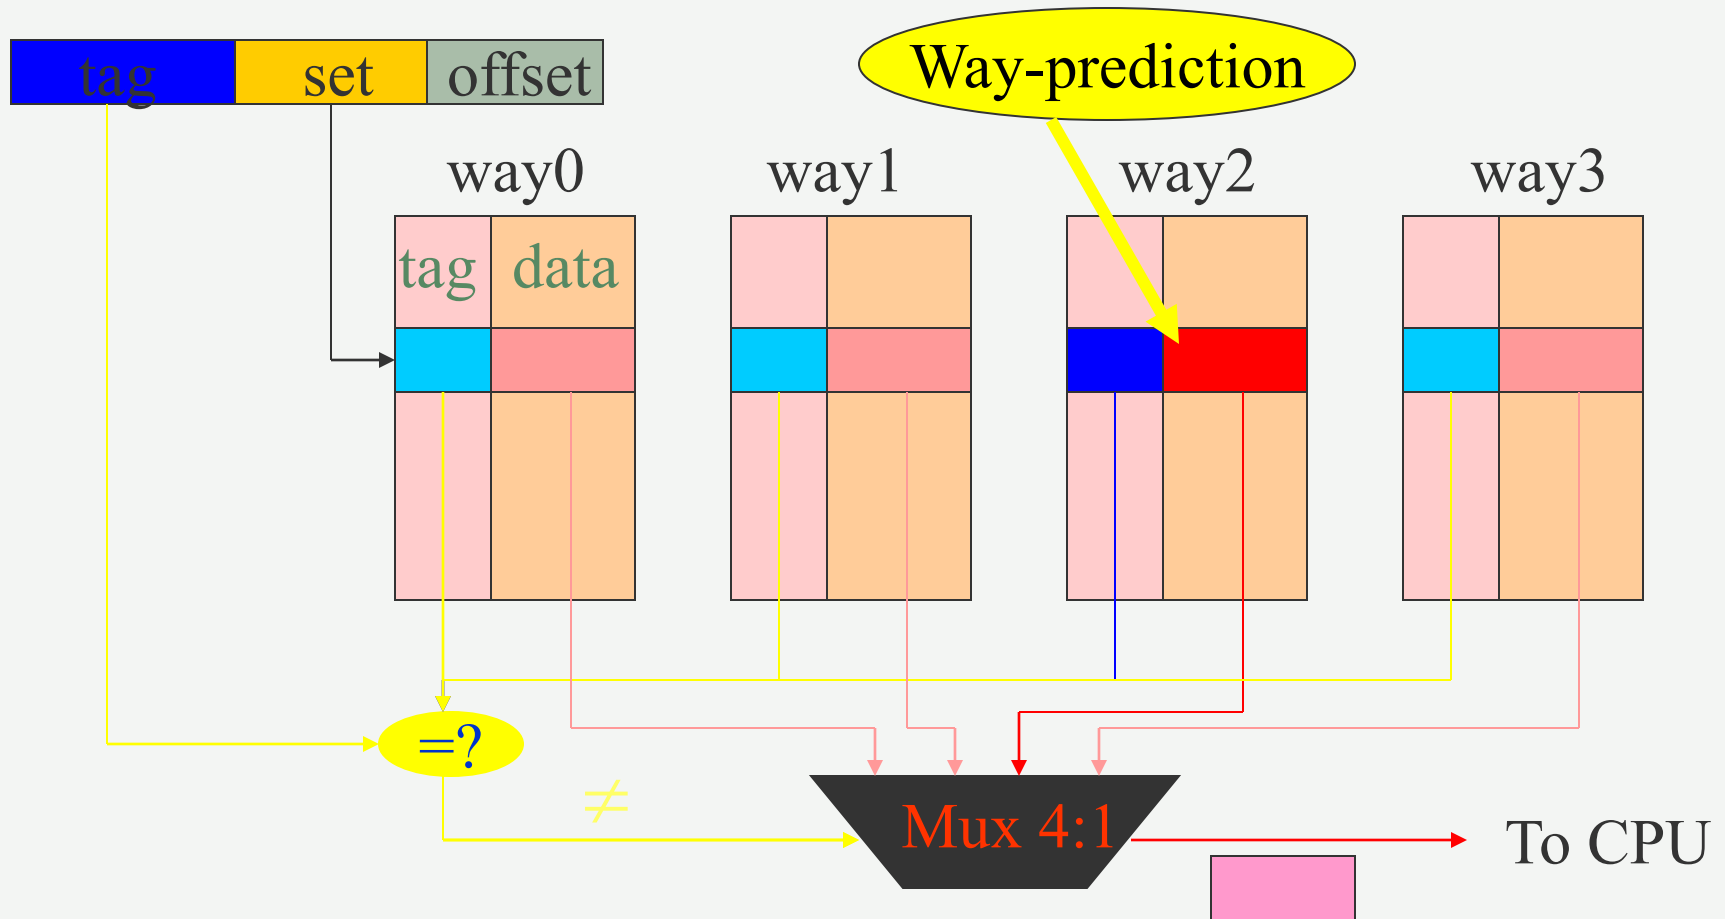
Best Case of Way Prediction: First Hit

Cost: way prediction only



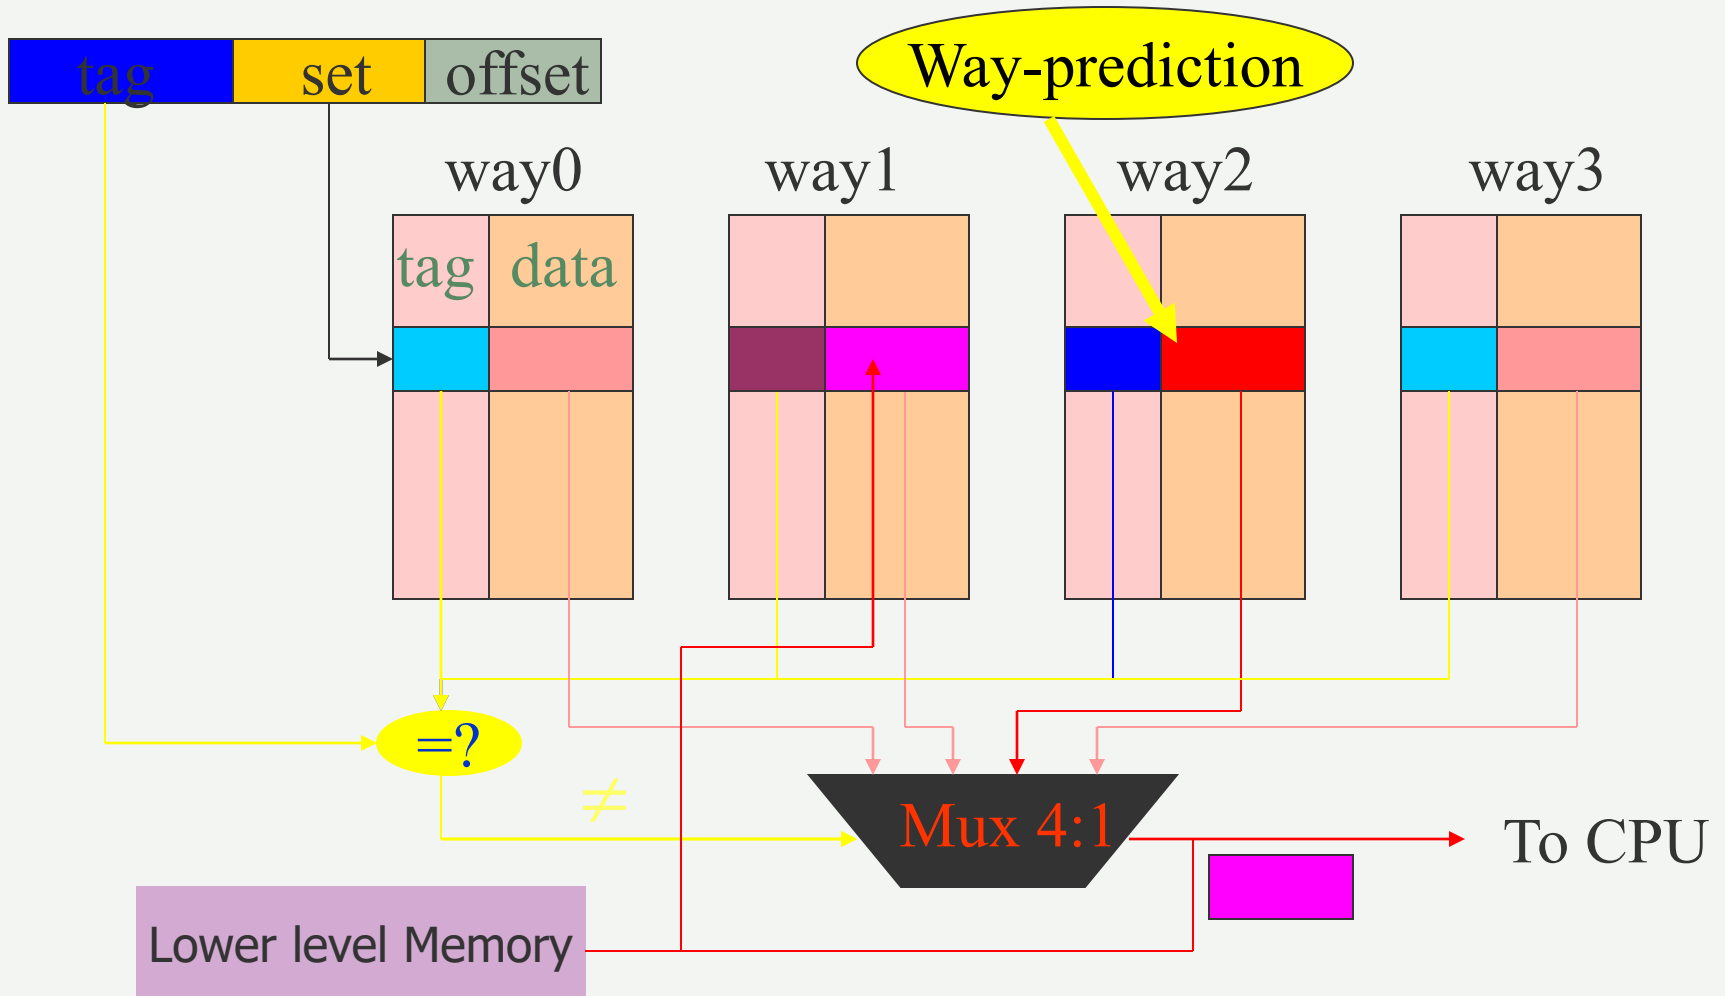
Way Prediction: Non-first Hit (in Set)

Cost: way prediction + selection in set



Worst Case of Way-prediction: Miss

Cost: way prediction + selection in set + miss



Make a **low-latency** and **high hit-rate** cache to be a reality

- A write to a set-associative cache

1. Map to the cache set
2. Find an empty way (block) to write
3. Or write to the least recent used (LRU) way (block)

- A read in a set-associative cache

1. Map to the cache set
2. Compare with the tags, then select the way if matches
3. Otherwise, serve the cache miss by going to memory

- Why a way-prediction is hard?

- **We do not know** which way the cache block is located
- Can we **enforce which way to go** in both write and read?

Basic Ideas

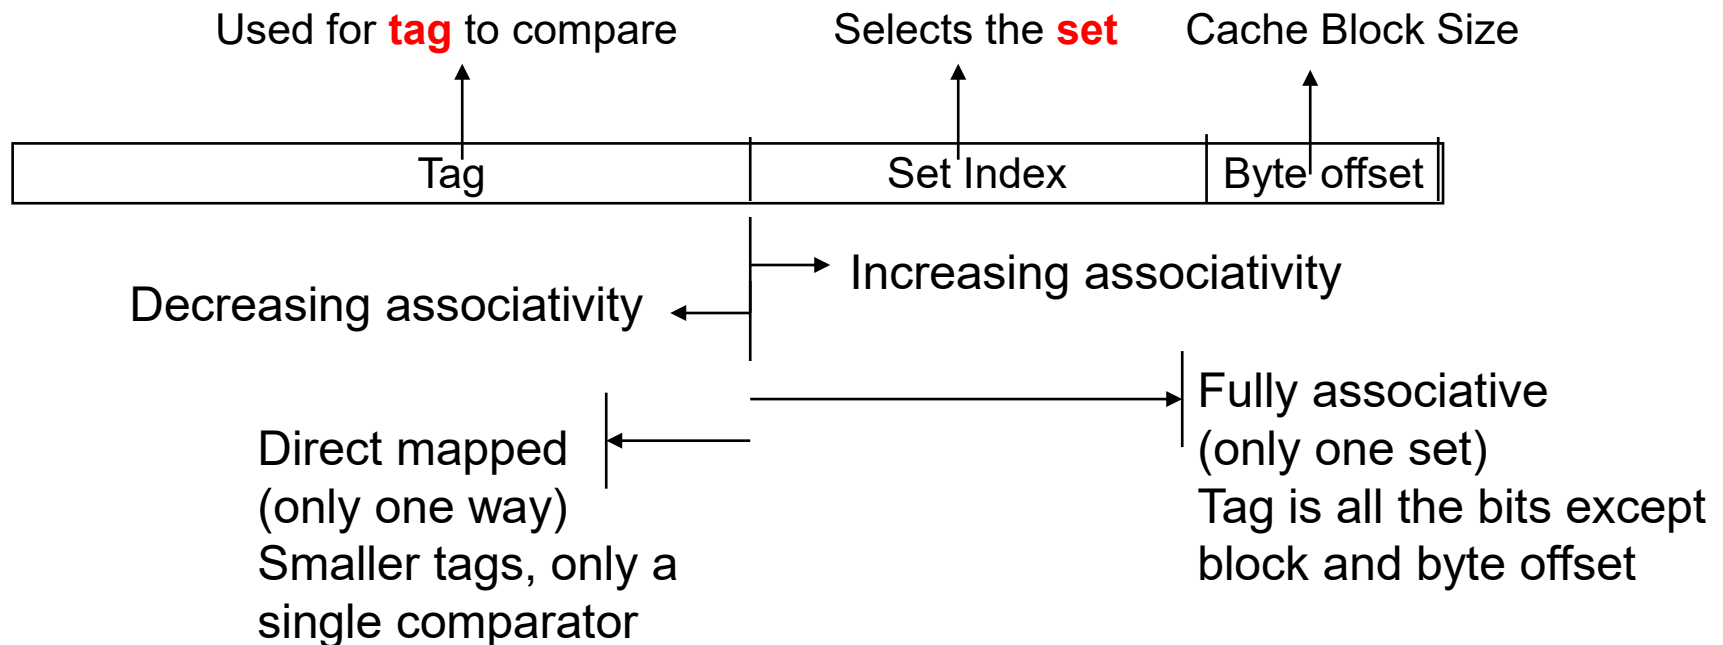
- Conventional Set Associative Cache
 - **Mapping** to a set of multiple ways
 - find a way (LRU) to **write**
 - search all ways for a **read**
 - **Latency** comes from “**finding**” or “**searching**” among **ways**
- Can we decisively or **directly** find a **way**?
 - If so, the search would be **direct** to **the way**
- How do we make this happen?
 - Multi-column cache

Multi-column Caches: Fast and High Hit Ratio


- Zhang, et. al., IEEE Micro, 1997.
- Objectives:
 - Each first hit is equivalent to a direct-mapped access.
 - Maximize the number of first hits.
 - Minimize the latency of non-first-hits.
 - Additional hardware should be simple and low cost.

Range of Set Associative Caches

- ❑ For a fixed size cache, each time the associativity increases it doubles the number of blocks per set (or the number of ways)
 - ❑ halves the number of sets – decreases the set index by 1 bit
 - ❑ increases the size of the tag by 1 bit
- ❑ Note: word offset is for multiple words in a block. word offset + byte offset = long cache block



Multi-Column Caches: Major Location

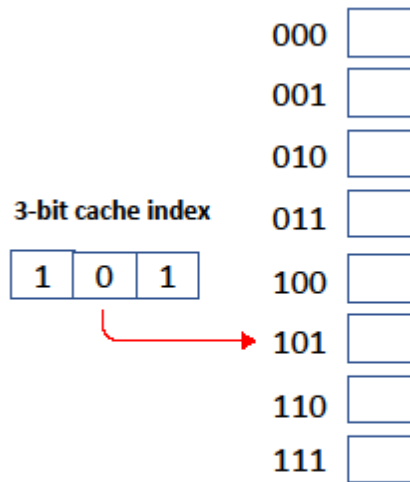


- The bits given to tag and ``set bits'' generate a direct-mapped location: **Major Location**.
- A major location mapping = **direct-mapping**.

Comparisons of Three Caches

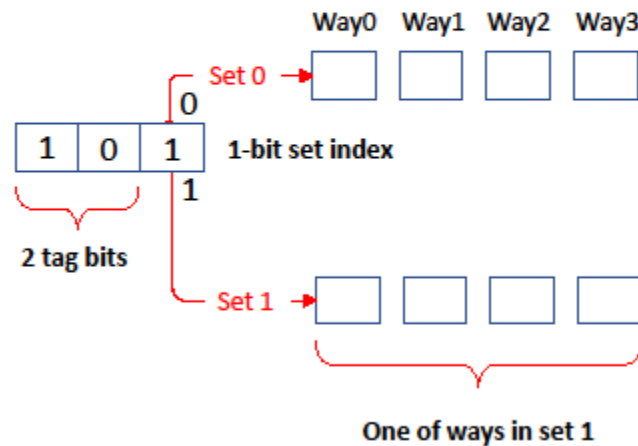
4-way Multicolumn Cache

Direct Map

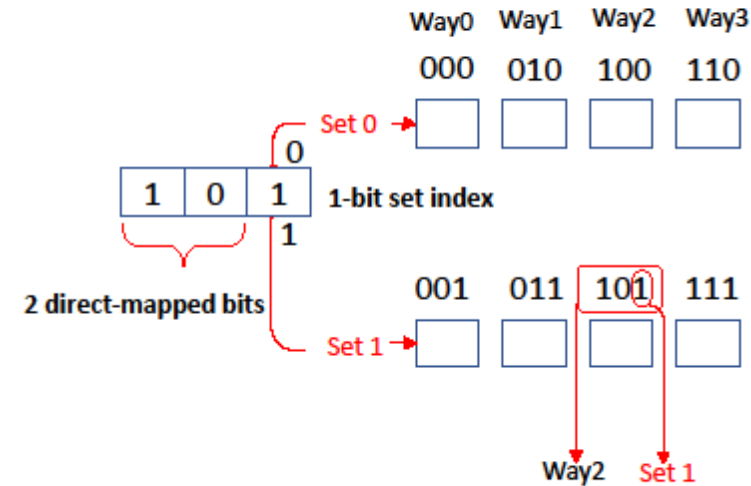


(a)

4-way Set Associative



(b)



(c)

Implicit direct mapping in set associative cache


- Standard set associative cache only maps to a set
- Multicolumn cache uses the full direct map bits to write and read in a set associative cache
 - **Set index** remains the same
 - The **bits given to the tag** from direct map are used for which **way** in the set
 - These two sets of bits are **available and free**
- Major location contains a **most recent used (MRU)** block either loaded from memory or just accessed.

Multi-column Caches: Selected Locations

- Multiple blocks can be direct-mapped to the same major location, but only MRU is the major.
- The non-MRU blocks are stored in other empty locations in the set: **Selected Locations**.
- If “other locations” are used for their **own** major locations, there will be no space for selected ones.
- **Swap:**
 - A block in **selected location** is **swapped** to **major location** as it becomes MRU.
 - A block in major location is **swapped** to a **selected location** after a new block is loaded in it from memory.

Multi-Column: Indexing Selected Locations

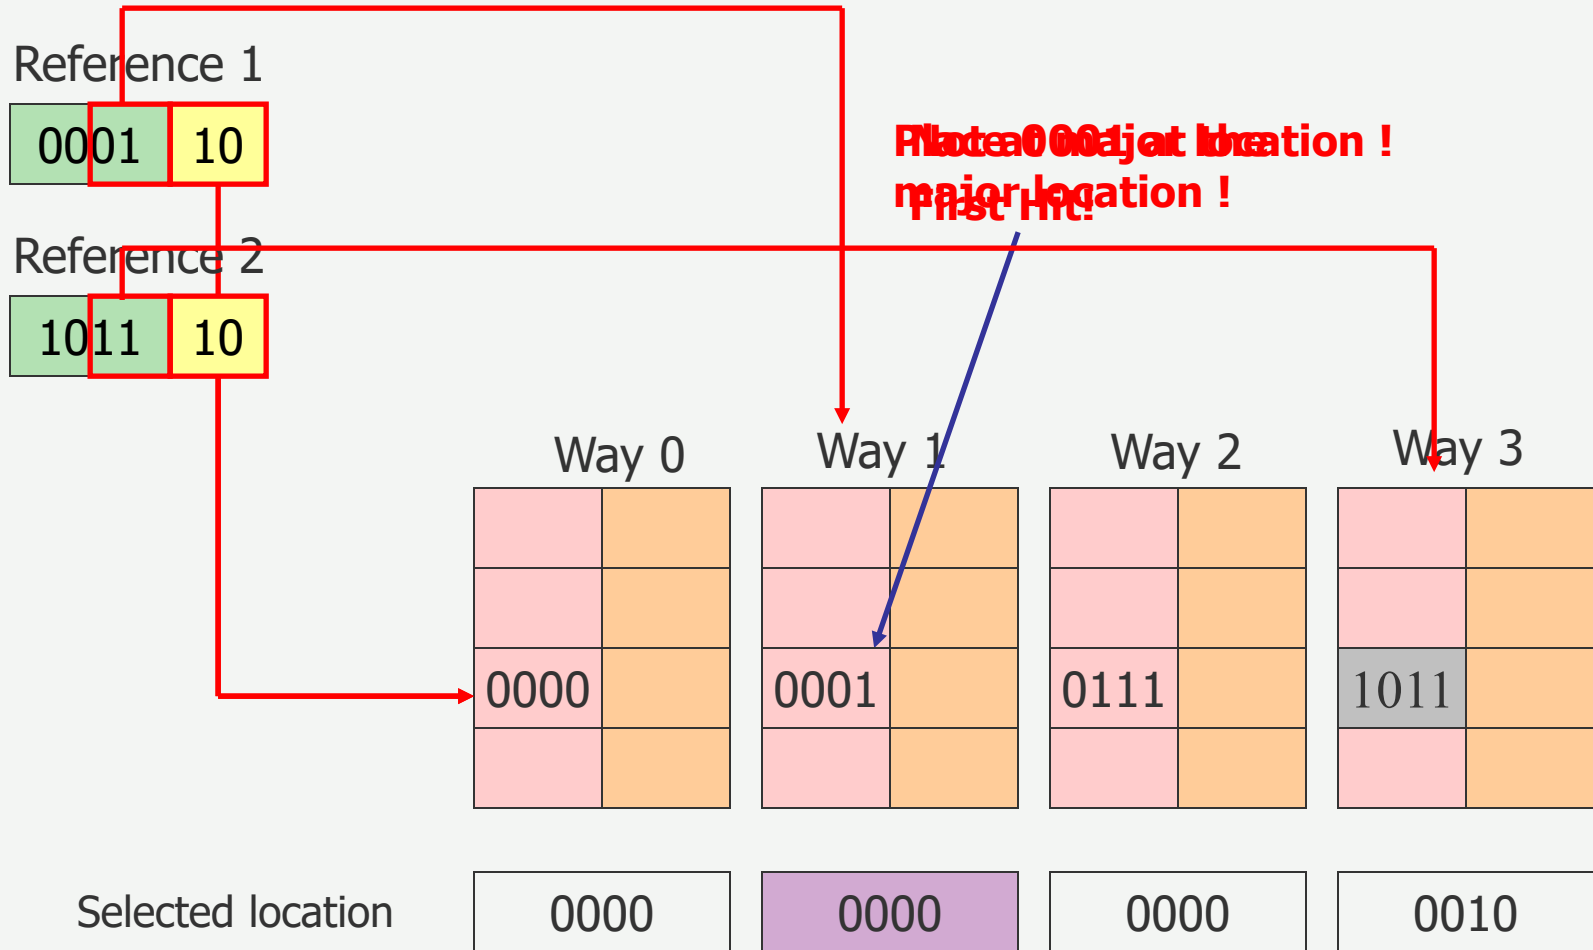
- The selected locations associated with its major location are indexed for a fast search.



Summary of Multi-column Caches

- A **major location** in a set is the direct-mapped location of MRU.
- A **selected location** is the direct-mapped location but non-MRU.
- A selected location **index** is maintained for each major location.
- A ``**swap**'' is used to ensure the block in the major location is always MRU.

Multi-column Cache Operations



Not a valid location !
Not a valid location !
First Hit!

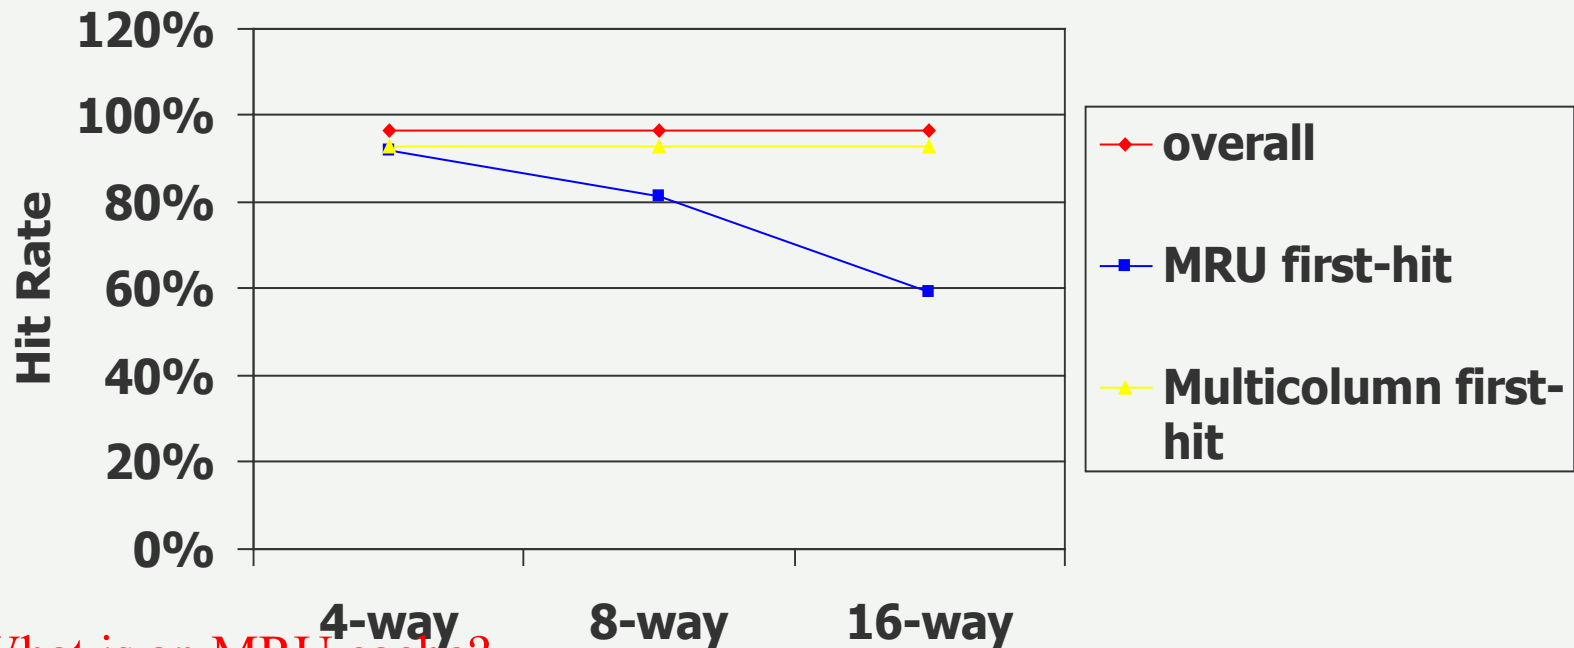
No selected location!

Performance Summary of Multi-column Caches

- Hit ratio to the **major locations** is about 90%.
- The hit ratio is higher than that of direct-mapped cache due to high associativity while keeps low access time of direct-mapped cache in average.
 - **First-hit** is equivalent to direct-mapped.
 - **Non-first hits** are faster than set-associative caches.
- Outperforming set associative caches

Comparing First-hit Ratios between Multicolumn and MRU Caches


64KB Data Cache Hit Rate (Program mgrid)



Note: What is an MRU cache?

An **MRU bit map** is used for a set associative cache, (1 bit for way, 2 bits for 4-way, ...), which remembers the most recent used way as the prediction of the next access. This bit map is used for the **way prediction**.

Multi-column Technique is Critical for Low Power Caches by targeting specific ways



Source: Intel.com

The impact of multicolumn caches

- The concept of multicolumn caches has influenced the R&D in the set-associative cache products:
 - A tag-less cache
 - K-way direct mapped cache
 - The direct-to-data cache
 - Highly associative caches for low power processors
 - A way memorization cache
 - Way selection based on tag bits
 - In **ARM Cortex R chip**, cache mapping “directly access one of the ways in a set”